



## MODULE 05 - Piscine Python for Data Science

Pandas: working with Dataframes

*Summary: Today we will help you acquire skills with Pandas.*

# Contents

<b>I</b>	<b>Foreword</b>	<b>2</b>
<b>II</b>	<b>Instructions</b>	<b>3</b>
<b>III</b>	<b>Specific instructions of the day</b>	<b>4</b>
<b>IV</b>	<b>Exercise 00 : Load and save</b>	<b>5</b>
<b>V</b>	<b>Exercise 01 : Basic operations</b>	<b>7</b>
<b>VI</b>	<b>Exercise 02 : Preprocessing</b>	<b>9</b>
<b>VII</b>	<b>Exercise 03 : Selects and aggregations</b>	<b>11</b>
<b>VIII</b>	<b>Exercise 04 : Enrichment and transformations</b>	<b>13</b>
<b>IX</b>	<b>Exercise 05 : Pandas optimizations</b>	<b>16</b>

# Chapter I

## Foreword

Fun facts about pandas:

- Pandas are big eaters – every day they fill their tummies for up to 12 hours, consuming up to 12 kilograms of bamboo
- Unlike most other bears, pandas do not hibernate. When winter approaches, they head lower down their mountain homes to warmer temperatures, where they continue to chomp away on bamboo
- Sadly, these beautiful bears are endangered, and it's estimated that only around 1,000 remain in the wild
- On average, pandas poo 40 times a day
- According to legend, the panda was once an all-white bear. When a small girl tried to save a panda cub from being attacked by a leopard, the leopard killed the girl instead. Pandas came to her funeral wearing armbands of black ashes. As they wiped their eyes, hugged each other, and covered the ears, they smudged the black ashes into their fur.
- A panda's entire mating process takes about two or three days. Once they have mated, the females chase the males out of their territory and raise their cubs on their own
- A giant panda usually gives birth to a single cub. Sometimes twins are born, but when this happens, the mother typically ignores the weaker cub. She does not have enough energy to care for two cubs
- A giant panda's face is cute, but it is not chubby. It gets its shape from massive cheek muscles
- Keeping even a single panda in a zoo is expensive. A panda costs five times more to keep than the next most expensive animal, an elephant
- If you think that any of this is relevant to the library Pandas, you needn't. The name is derived from the term "panel data", an econometrics term for data sets that include observations over multiple time periods for the same individuals

# Chapter II

## Instructions

- Use this page as your only reference. Do not listen to any rumors or speculations about how to prepare your solution.
- Here and further on we use Python 3 as the only correct version of Python.
- The solutions for python exercises (d01, d02, d03) must have the following block in the end: `if __name__ == '__main__':`
- Pay attention to the permissions of your files and directories.
- To be assessed your solution must be in your GIT repository.
- Your solutions will be evaluated by your piscine peers.
- You should not leave any additional files in your directory other than those explicitly specified in the subject. It is recommended that you modify your .gitignore to avoid any accidents.
- When you need to get precise output in your programs, it is forbidden to display a precalculated output instead of performing the exercise correctly.
- Have a question? Ask your neighbor on the right. If that fails, try your neighbor on the left.
- Your reference material: peers / Internet / Google.
- You can ask questions in Slack.
- Read the examples carefully. They may require things that are not otherwise specified in the subject.
- And may the Force be with you!


# Chapter III

## Specific instructions of the day

- Use Jupyter Notebook to work with your code
- For each major subtask in the list of any exercise (black bullets), your ipynb file should have an h2 heading to help your peer easily navigate in your code
- No imports are allowed, except those explicitly mentioned in the “Authorized functions” section of the title block of each exercise
- You can use any built-in function, if it is not prohibited in the exercise
- Save and load all the required data in the subfolder data/

# Chapter IV

## Exercise 00 : Load and save

	Exercise 00
Load and save	
Turn-in directory : <i>ex00/</i>	
Files to turn in : <code>load_and_save.ipynb</code>	
Allowed functions : <code>import pandas as pd</code>	

Congratulations! Five days and one rush are behind you. They were dedicated to building good fundamental skills in working with Python. Now you know all the basic data types, useful built-in functions, how to work with a virtual environment, how to use OOP, how to use logging, how to parse data from websites, and how to send messages in a Slack channel. In the following week, we will be more focused on data science. It will be a week full of Pandas exercises. Pandas is one of the most popular and useful libraries in the field. You will love it too!

In this exercise, you will need to load the [log file](#) (put it in the directory data in the root directory of the day) into a dataframe, change the delimiter, and save it to another file.

The task is:

- `read_csv`:
  - filter the rows with index 2 and 3 using the argument `skiprows`, we know that these observations were fake
  - filter the last 2 rows from the footer using the argument `skipfooter`, we know that these observations were fake too
  - assign the following names to the column: `datetime`, `user`
  - use `datetime` as the index column
- rename `datetime` to `date_time`
- `to_csv`:
  - use `';` as the delimiter

- save it to a file with the name feed-views-semicolon.log

As the result of `read_csv`, you need to achieve the following:


```
In [3]: df.head()
Out[3]:
      datetime      user
2020-04-17 12:01:08.463179 artem
2020-04-17 12:01:23.743946 artem
2020-04-17 12:35:52.735016 artem
2020-04-17 12:36:21.401412 oksana
2020-04-17 12:36:22.023355 oksana

In [4]: df.tail()
Out[4]:
      datetime      user
2020-05-21 16:36:40.915488 ekaterina
2020-05-21 17:49:36.429237 maxim
2020-05-21 18:45:20.441142 valentina
2020-05-21 23:03:06.457819 maxim
2020-05-21 23:23:49.995349 pavel
```

Think about how many lines of code you would have needed to write if you had had to do it without Pandas

# Chapter V

## Exercise 01 : Basic operations

	Exercise 01
	Basic operations
	Turn-in directory : <i>ex01/</i>
	Files to turn in : <code>basic_operations.ipynb</code>
	Allowed functions : <code>import pandas as pd</code>

We are confident that you understand that this is not everything that Pandas can do for you. Let us go deeper and wider.

In this exercise, you will work with a single log of users who visited a page, including their timestamps.


- create a dataframe `views` with two columns: `datetime` and `user` by reading `feed-views.log`
  - convert the `datetime` to the `datetime64[ns]` Dtype
  - extract the year, month, day, hour, minute, and second from the values of that column to the new columns
- create the new column `daytime`
  - you need to assign the particular time of day value if an hour is within a particular interval, for example, `afternoon` if the hour is larger than 11 and less or equal to 17
  - 0.00 – 03.59 `night`, 04.00 – 06.59 `early morning`, 07.00 – 10.59 `morning`, 11.00 – 16.59 `afternoon`, 17.00 – 19.59 `early evening`, 20.00 – 23.59 `evening`
  - use the method `cut` to solve this subtask
  - assign the column `user` as the index
- calculate the number of elements in your dataframe
  - use the method `count()`



- calculate the number of elements in each time of day category using the method `value_counts()`
- sort values in your dataframe by hour, minute, and second in ascending order (simultaneously and not one by one)
- calculate the minimum and maximum for the hours and the mode for the daytime categories
  - calculate the maximum of hour for the rows where the time of day is night
  - calculate the minimum of hour for the rows where the time of day is morning
  - In addition to this, find out who visited the page at those hours (make one example from that)
  - calculate the mode for the hour and daytime
- show the 3 earliest hours in the morning and the corresponding usernames and the 3 latest hours and the usernames using `nsmallest()` and `nlargest()`
- use the method `describe()` to get the basic statistics for the columns
  - to find out what the most popular interval for visiting the page is, calculate the interquartile range for the hour by extracting values from the result of the `describe()` method and store it in the variable `iqr`

# Chapter VI

## Exercise 02 : Preprocessing

	Exercise 02
Preprocessing	
Turn-in directory : <i>ex02/</i>	
Files to turn in : <b>preprocessing.ipynb</b>	
Allowed functions : <b>import pandas as pd</b>	

One day you will train machine learning models (no later than this week), but most of them require the data to be clean and enriched: without duplicates or missing values. Pandas is a great tool to do this. It gives you the tools not only to perform descriptive analysis and understand your data better but to preprocess it too. That is what you are going to do in this exercise.


- [download](#) and read the CSV file and make ID the index column
- count the number of observations using the method `count()`
- drop the duplicates, taking into account only the following columns: CarNumber, Make\_n\_model, Fines
  - between the two equal observations, you need to choose the last
  - check the number of observations again
- work with missing values
  - check how many values are missing from each column
  - drop all the columns with over 500 missing values using the argument `thresh`, check how many missing values are in each column
  - replace all the missing values in the Refund column with the previous value in that column for that cell, use the argument `method`, check how many values are missing from each column
  - replace all the missing values in the Fines column with the mean value of this column (exclude NA/null values when computing the mean value), check how many values are missing from each column

- split and parse the make and model
  - use the method `apply` both for splitting and for extracting the values to the new columns `Make` and `Model`
  - drop the column `Make_model`
  - save the dataframe in the JSON file `auto.json` in the format below:

```
[{"CarNumber": "Y16308161RUS", "Refund": 2.0, "Fines": 3200.0, "Make": "Ford", "Model": "Focus"}, {"CarNumber": "E432XX77RUS", "Refund": 1.0, "Fines": 6500.0, "Make": "Toyota", "Model": "Camry"}]
```

# Chapter VII

## Exercise 03 : Selects and aggregations

	Exercise 03
Selects and aggregations	
Turn-in directory : <i>ex03/</i>	
Files to turn in : <b>selects_n_aggs.ipynb</b>	
Allowed functions : <b>import pandas as pd</b>	


Ok, great. Now we have cleaned our data: all the duplicates have been dropped, all the missing values have been deleted, and our columns have been reorganized to be more convenient for analysis. Now we can go further. We have a lot of questions that need to be answered.

- load the JSON file that you created in the previous exercise into a dataframe
  - set CarNumber as the index column
- make the following selects
  - display the rows only where the fines are more than 2,100
  - display the rows only where the fines are more than 2,100 and the refund equals 2
  - display the rows only where the models are from the list: ['Focus', 'Corolla']
  - display the rows only where the car number is from the list: ['Y7689C197RUS', '92928M178RUS', '7788KT197RUS', 'H115YO163RUS', 'X758HY197RUS']
- make the aggregations with the make and the model
  - display the median fines grouped by the make
  - display the median fines grouped by the make and the model

- display the number of fines grouped by the make and the model in order to understand if we can trust the median values
- display the minimum and the maximum fines grouped by the make and the model in order to better understand the variance
- display the standard deviation of the fines grouped by the make and the model in order to better understand the variance
- make the aggregations with the car number
  - display the car numbers grouped by the number of the fines in descending order, we want to find those who most often violated the law
  - select from the initial dataframe all the rows corresponding to the top-1 car number, we want to zoom in a little bit
  - display the car numbers grouped by the sum of the fines in descending order, we want to find those who paid the most
  - select from the initial dataframe all the rows corresponding to the top-1 car number, we want to zoom in a little bit
  - display a table that answers the question: are there any car numbers that were connected to different models?

# Chapter VIII

## Exercise 04 : Enrichment and transformations

	Exercise 04
Enrichment and transformations	
Turn-in directory : <i>ex04/</i>	
Files to turn in : <code>enrichment.ipynb</code>	
Allowed functions : <code>import pandas as pd, import numpy as np, import requests</code>	

Cool. But the more data you have the better the analysis you can conduct. Let us enrich our initial dataset.

- read the JSON file that you saved in `ex02`
  - one of the columns has the float type, so let us define the format of it in pandas using `pd.options.display.float_format`: floats should be displayed with two decimals
  - there are values missing from the Model, do not do anything with them
- enrich the dataframe using a sample from that dataframe
  - create a sample with 200 new observations with `random_state = 21`
    - \* the sample should not have new combinations of the car number, make and model, so the whole dataset will be consistent in these terms
    - \* there are no restrictions on the refund and fines, you can take any value from these columns at random and use it towards any car number
  - concatenate the sample with the initial dataframe to a new dataframe `concat_rows`
- enrich the dataframe `concat_rows` by a new column with the data generated
  - create a series with the name `Year` using random integers from 1980 to 2019

- use `np.random.seed(21)` before generating the years
- concatenate the series with the dataframe and name it `fines`
- enrich the dataframe with the data from another dataframe
  - create a new dataframe with the car numbers and their owners
    - \* get the most popular surnames(**you can find the file `surname.json` in the attachments**) in the US
    - \* create a new series with the surnames (they should not have special characters like commas, brackets, etc.) from the data you gathered, the count should be equal to the number of unique car numbers using the sample (use `random_state = 21`)
    - \* create the dataframe `owners` with 2 columns: `CarNumber` and `SURNAME`
  - append 5 more observations to the `fines` dataframe (come up with your own ideas of `CarNumber`, etc.)
  - delete the dataframe last 20 observations from the `owners` and add 3 new observations (they are not the same as those you add to the `fines` dataframe)
  - join both dataframes:
    - \* the new dataframe should have only the car numbers that exist in both dataframes
    - \* the new dataframe should have all the car numbers that exist in both dataframes
    - \* the new dataframe should have only the car numbers from the `fines` dataframe
    - \* the new dataframe should have only the car numbers from the `owners` dataframe
- create a pivot table from the `fines` dataframe, it should look like this (the values are the sums of the fines), but with all the years (the values may be different for you):


	Year	1980	1981	1982	1983	1984		
Make	Model							
Dodge	Neon	nan	nan	nan	nan	nan		
Ford	Focus	89,194.59		266,783.76	107,283.76	147,289.17	106,000.00	
	Mondeo	nan	nan	46,200.00	nan	nan		
Honda	CRV	nan	nan	nan	nan	nan		
Hyundai	Solaris	nan	nan	nan	nan	nan		
Lada	Kalina	nan	nan	nan	nan	nan		
Merc	S300	nan	nan	nan	nan	nan		
Skoda	Octavia	13,794.59		1,900.00	8,894.59	nan	1,300.00	
Toyota	Camry	12,000.00		nan	1,000.00	8,594.59	1,000.00	
	Corolla	nan	6,800.00		nan	12,800.00	nan	
Volkswagen	Golf		20,800.00		8,594.59	5,000.00	200.00	nan
	Jetta	nan	1,000.00		nan	nan	nan	
	Passat	900.00	12,500.00		nan	1,100.00	8,594.59	
	Touareg	nan	nan	nan	nan	nan		

- save both the fines and owners dataframes to CSV files without an index



# Chapter IX

## Exercise 05 : Pandas optimizations

	Exercise 05
Pandas optimizations	
Turn-in directory : <i>ex05/</i>	
Files to turn in : <code>optimizations.ipynb</code>	
Allowed functions : <code>import pandas as pd, import gc</code>	

We are returning to the idea of code efficiency. By now you know the basics of Pandas. It is time to get to know some cool stuff that most Pandas users neither use nor know about.

- read the `fines.csv` that you saved in the previous exercise
- iterations: in all the following subtasks, you need to calculate `fines/refund*year` for each row and create a new column with the calculated data and measure the time using the magic command `%%timeit` in the cell
  - loop: write a function that iterates through the dataframe using for `i` in `range(0, len(df))`, `iloc` and `append()` to a list, assign the result of the function to a new column in the dataframe
  - do it using `iterrows()`
  - do it using `apply()` and lambda function
  - do it using Series objects from the dataframe
  - do it as in the previous subtask but with the method `.values`
- indexing: measure the time using the magic command `%%timeit` in the cell
  - get a row for a specific CarNumber, for example, `'O136HO197RUS'`
  - set the index in your dataframe with CarNumber
  - again, get a row for the same CarNumber

- downcasting:
  - run `df.info(memory_usage='deep')`, pay attention to the Dtype and the memory usage
  - make a `copy()` of your initial dataframe into another dataframe optimized
  - downcast from float64 to float32 for all the columns
  - downcast from int64 to the smallest numerical dtype possible
  - run `info(memory_usage='deep')` for your new dataframe, pay attention to the Dtype and the memory usage
- categories:
  - change the object type columns to the type category
  - This time, check the memory usage, it probably has a decrease of 2-3 times compared to the initial dataframe
- memory clean
  - using `%reset_selective` and the library `gc` clean the memory of your initial dataframe only